



SYNTHEX EVIDENCE REPORT

Executive Summary · for Decision makers

What this report proves

Synthex turned web data into court-grade evidence: every source was fetched, screened for malicious content before reaching the LLM, classified, and cryptographically sealed. The seal proves WHEN and that the bytes are UNCHANGED — not that the underlying claims are true.

Target Competitor X

Lens gtm

Fetched

2026-05-30T12:09:39.210Z

Sealed 2026-05-30T12:09:39.685Z

Sources 5

Blocked pre-LLM (FORGE) 0

Cryptographic seal

```
SHA-256      :  
c605b0572cc4c5bd00c37311408ae1a30fa5f83f177700760721d4dbef1bf616  
HMAC-SHA256  :  
ffd7884ebb244b4c5fc89f5d4f520a4fae86087f1a14deb0868fbfa43a74a415  
Seal method  : HMAC-SHA256 + RFC 3161 TSA  
RFC 3161 TSA : digicert (RFC 3161)  
  genTime    : 2026-05-30T12:09:39.000Z  
  serial     : 00fe0492cbce074002e7c6012387ecda8a
```



Scan the QR or follow page 6 to verify this seal yourself, offline, with OpenSSL.



Threats blocked by FORGE (pre-LLM)

FORGE is a deterministic regex pre-filter that runs BEFORE any LLM call. It blocks injection and exfiltration vectors so they never reach (or poison) the model. Categories are mapped to their OWASP reference class. The pre-filter is heuristic, not formally verified.

No malicious content detected in this batch.

All fetched sources passed the FORGE pre-filter and were forwarded to classification.

Classification severity (per lens)

[GTM] severity 7/10

Recorte de precio agresivo (-20%) + free tier: presión competitiva alta.

signals: price-cut · free-tier · 10k-api-calls

[GTM] severity 6/10

Expansión de ventas enterprise en EMEA: señal de go-to-market.

signals: hiring · enterprise-sales · EMEA

[GTM] severity 6/10

Documentación de prevención de prompt-injection (informativa): describe la técnica, no la ejecuta.

signals: prompt-injection-docs · owasp · mitigations



SYNTHEX EVIDENCE REPORT

Cost & Efficiency · for CFO

Token cost saved (deduplication)

FORGE deduplicates fetched documents by full SHA-256 before classification: every duplicate block removed is one fewer LLM call paid for. The ratio below is exact (not estimated).

Unique blocks 4





Duplicate blocks 1

Dedup ratio 20.0% (LLM calls avoided)

Bytes saved 93 bytes not sent to the model

Latency per stage (wall-clock)

Measured by the pipeline (evidence.timings), outside the sealed payload so it never affects verification.

FETCH		0.5 ms
FORGE		18.3 ms
ALIGNMENT_CHECK		0.3 ms
CLASSIFY		0.2 ms
PROVE		472.3 ms
		Total: 491.6 ms



EU AI Act — Article 12 (record-keeping / logging)

Article 12 requires high-risk AI systems to automatically record events ('logs') enabling traceability. This report is itself such a record: an immutable, timestamped log of the data lifecycle, sealed so any tampering is detectable.

OK	Traceability of inputs	5 source URL(s) recorded
OK	Pre-processing log	0 block(s) by FORGE pre-filter, with category
OK	Classification output	3 finding record(s) retained
OK	Integrity of records	HMAC-SHA256 over canonical payload (present)
OK	Timekeeping (independent)	RFC 3161 TSA — digicert, 2026-05-30T12:09:39.000Z
OK	Tamper-evidence	SHA-256 content hash: c605b0572cc4c5bd00c37311...
OK	Record timestamp	2026-05-30T12:09:39.685Z

This is an evidentiary record, not legal advice. Mapping to Article 12 is provided as a good-faith aid. The deploying system is classified high-risk.



SYNTHEX EVIDENCE REPORT

Risk Snapshot · for Broker / Underwriter

Synthex Risk Score (CVSS 0–10 severity scale · NIST AI RMF / EU AI Act framing · mapping, not endorsement)

49

/ 100

MEDIUM RISK

How this number is computed (honest formula)

```
maxSev    = max severity across all findings ..... 7/10
blockTerm  = min(blockedCount, 5) / 5 * 10 ..... 0.0/10 (blocked=0)
score      = round( (maxSev*0.70 + blockTerm*0.30) * 10 ) = 49
```

DISCLAIMER

This Risk Score is computed by Synthex from the data in this report using the deterministic formula shown above. Its severity axis is the CVSS 0–10 base-score scale and its bands align to CVSS severity ratings; the compliance framing maps to NIST AI RMF and EU AI Act risk categories. This is a MAPPING, NOT an ENDORSEMENT: it is NOT a Munich Re assessment, NOT an insurance rating, and NOT underwriting advice. No third party or framework body has reviewed or endorsed this number. See docs/compliance-mapping.md.

Verdict

HIGH RISK — max severity 7/10 across 3 classified source(s); lead signal: price-cut.

3 questions this evidence raises

1. Which of the 3 classified source(s) most changes our gtm exposure, and who owns the response?
2. Do the surfaced signals cross our escalation threshold under the gtm policy?
3. Is the sealed evidence (hash + timestamp + signature) sufficient to defend the highest-severity finding in an audit?



Reproduce the proof offline

RFC 3161 proves WHEN this content was sealed — not that the content is TRUE. The HMAC proves integrity against the issuing key. Reproduce the checks yourself with OpenSSL:

```
# 1. Decode the trusted timestamp token shipped in the report JSON (seal.rfc3161Tsa.token):  
echo "<base64-token>" | base64 -d > synthex.tsr  
  
# 2. Inspect the DigiCert RFC 3161 reply (works offline):  
openssl ts -reply -in synthex.tsr -text  
  
# 3. Confirm the messageImprint in that output equals this SHA-256:  
#    c605b0572cc4c5bd00c37311408ae1a30fa5f83f177700760721d4dbef1bf616  
  
# 4. Full chain validation (needs the DigiCert CA chain):  
openssl ts -verify -in synthex.tsr -CAfile digicert-chain.pem -sha256
```

```
c605b0572cc4c5bd00c37311408ae1a30fa5f83f177700760721d4dbef1bf616
```